

Operating System

MP1: System Call

Team Member: 林子謙、陳奕潔

Contributions: Both contributed to system call implementing and report writing.

Contents

1	Code Tracing	3
1.1	SC_Halt	3
1.1.1	Machine::Run()	3
1.1.2	Machine::OneInstruction()	3
1.1.3	Machine::RaiseException()	5
1.1.4	ExceptionHandler()	5
1.1.5	SysHalt()	7
1.1.6	Interrupt::Halt()	7
1.2	SC_Create	7
1.2.1	ExceptionHandler()	7
1.2.2	SysCreate()	9
1.2.3	FileSystem::Create()	9
1.3	SC_PrintInt	9
1.3.1	ExceptionHandler()	9
1.3.2	SysPrint()	11
1.3.3	SynchConsoleOutput::PutInt()	12
1.3.4	ConsoleOutput::PutChar()	12
1.3.5	Interrupt::Schedule()	13
1.3.6	Interrupt::CheckIfDue()	13
1.3.7	ConsoleOutput::CallBack()	15
1.3.8	SynchConsoleOutput::CallBack()	16
1.3.9	Machine::Run()	16
1.3.10	Interrupt::OneTick()	17
1.4	Makefile	17
1.4.1	Variable Definitions	19
1.4.2	Shell Commands and Dependencies	19
1.4.3	Clean-Up and Error Messages	20
2	System Call Implementation	20
2.1	Open	23
2.1.1	ExceptionHandler()	23
2.1.2	SysOpen()	23
2.1.3	OpenAFile()	24
2.2	Write	25
2.2.1	ExceptionHandler()	25
2.2.2	SysWrite()	25
2.2.3	WriteFile()	26
2.3	Read	26
2.3.1	ExceptionHandler()	26

2.3.2	SysRead()	26
2.3.3	ReadFile()	26
2.4	Close	28
2.4.1	ExceptionHandler()	28
2.4.2	SysClose()	28
2.4.3	CloseFile()	29
3	Difficulties	29

1 Code Tracing

1.1 SC_Halt

The SC_Halt system call is designed to shut down NachOS.

1.1.1 Machine::Run()

```
54 void Machine::Run() {
55     Instruction *instr = new Instruction; // storage for decoded instruction
56     if (debug->IsEnabled('m')) {
57         cout << "Starting program in thread: " << kernel->currentThread->getName();
58         cout << ", at time: " << kernel->stats->totalTicks << "\n";
59     }
60     kernel->interrupt->setStatus(UserMode);
61     for (;;) {
62         DEBUG(dbgTraCode, "In Machine::Run(), into OneInstruction "
63             | | | | | << "==" Tick " << kernel->stats->totalTicks << " ==");
64         OneInstruction(instr);
65         DEBUG(dbgTraCode, "In Machine::Run(), return from OneInstruction "
66             | | | | | << "==" Tick " << kernel->stats->totalTicks << " ==");
67
68         DEBUG(dbgTraCode, "In Machine::Run(), into OneTick "
69             | | | | | << "==" Tick " << kernel->stats->totalTicks << " ==");
70         kernel->interrupt->OneTick();
71         DEBUG(dbgTraCode, "In Machine::Run(), return from OneTick "
72             | | | | | << "==" Tick " << kernel->stats->totalTicks << " ==");
73         if (singleStep && (runUntilTime <= kernel->stats->totalTicks))
74             Debugger();
75     }
76 }
```

Figure 1: Machine::Run()

Machine::Run() method initializes the status to **User Mode** and then enters an infinite loop. Within the loop, it **fetches an instruction** by calling OneInstruction(instr) on line 64 and **increments the system time** by calling kernel->interrupt->OneTick(). You may refer to Figure 1.

1.1.2 Machine::OneInstruction()

Once Machine::OneInstruction() is invoked by Machine::Run(), it reads an instruction from the program counter and decodes it (on lines 133 to 136, you may refer to Figure 2).

Next, the instruction is handled by going through several switch cases based on its opcode. In the **Sys_Halt** example, Machine::OneInstruction() enters the **OP_SYSCALL** case (lines 663 to 666, see Figure 3), invokes RaiseException(), and passes the corresponding exception type to it.

```

122 void Machine::OneInstruction(Instruction *instr) {
123     #ifdef SIM_FIX
124     |     int byte; // described in Kane for LWL,LWR,...
125     #endif
126
127     int raw;
128     int nextLoadReg = 0;
129     int nextLoadValue = 0; // record delayed load operation, to apply
130     | | | | | // in the future
131
132     // Fetch instruction
133     if (!ReadMem(registers[PCReg], 4, &raw))
134     |     return; // exception occurred
135     instr->value = raw;
136     instr->Decode();
137
138     if (debug->IsEnabled('m')) {
139     |     struct OpString *str = &opStrings[instr->opCode];
140     |     char buf[80];
141
142     |     ASSERT(instr->opCode <= MaxOpcode);
143     |     cout << "At PC = " << registers[PCReg];
144     |     sprintf(buf, str->format, TypeToReg(str->args[0], instr),
145     |         |     TypeToReg(str->args[1], instr), TypeToReg(str->args[2], instr));
146     |     cout << "\t" << buf << "\n";
147     | }
148
149     // Compute next pc, but don't install in case there's an error or branch.
150     int pcAfter = registers[NextPCReg] + 4;
151     int sum, diff, tmp, value;
152     unsigned int rs, rt, imm;
153
154     // Execute the instruction (cf. Kane's book)
155 > switch (instr->opCode) {...
684
685     // Now we have successfully executed the instruction.
686
687     // Do any delayed load operation
688     DelayedLoad(nextLoadReg, nextLoadValue);
689
690     // Advance program counters.
691     registers[PrevPCReg] = registers[PCReg]; // for debugging, in case we
692     | | | | | | | | | // are jumping into lala-land
693     registers[PCReg] = registers[NextPCReg];
694     registers[NextPCReg] = pcAfter;
695 }

```

Figure 2: Machine::OneInstruction()

```

663 |         case OP_SYSCALL:
664 |             DEBUG(dbgTrnCode, "In Machine::OneInstruction, RaiseException(SyscallException, 0), " << kernel->stats->totalTicks);
665 |             RaiseException(SyscallException, 0);
666 |             return;
667 |
668 |         case OP_XOR:
669 |             registers[instr->rd] = registers[instr->rs] ^ registers[instr->rt];
670 |             break;
671 |
672 |         case OP_XORI:
673 |             registers[instr->rt] = registers[instr->rs] ^ (instr->extra & 0xffff);
674 |             break;
675 |
676 |         case OP_RES:
677 |         case OP_UNIMP:
678 |             RaiseException(illegalInstrException, 0);
679 |             return;
680 |
681 |         default:
682 |             ASSERT(FALSE);
683 |     }

```

Figure 3: Machine::OneInstruction() case: OP_SYSCALL

1.1.3 Machine::RaiseException()

Since this instruction is a type of **system call**, it must be handled in **kernel mode**. Therefore, the system invokes Machine::RaiseException() to **switch to kernel mode** and clean up all executed instructions by calling DelayedLoad(0, 0);, and then calls ExceptionHandler() to handle the exception. After returning from ExceptionHandler(), the system may **switch back to user mode**. You may refer to Figure 4 for details.

```

97 | void Machine::RaiseException(ExceptionType which, int badVAddr) {
98 |     DEBUG(dbgMach, "Exception: " << exceptionNames[which]);
99 |     registers[BadVAddrReg] = badVAddr;
100 |     DelayedLoad(0, 0); // finish anything in progress
101 |     kernel->interrupt->setStatus(SystemMode);
102 |     ExceptionHandler(which); // interrupts are enabled at this point
103 |     kernel->interrupt->setStatus(UserMode);
104 | }

```

Figure 4: Machine::RaiseException()

1.1.4 ExceptionHandler()

ExceptionHandler() first **retrieves the system call exception type from register r2**, as defined in userprog/syscall.h (see Figure 5, line 53). In the case of Halt(), the system call exception type code is defined as 0.

Next, based on the system call exception type code, the program enters the SC_Halt case and calls SysHalt(). (As shown in Figure 6.)

```

50 void ExceptionHandler(ExceptionType which) {
51     char ch;
52     int val;
53     int type = kernel->machine->ReadRegister(2);
54     int status, exit, threadID, programID, fileID, numChar;
55     DEBUG(dbgSys, "Received Exception " << which << " type: " << type << "\n");
56     DEBUG(dbgTraCode, "In ExceptionHandler(), Received Exception "
57         << which << " type: " << type << ", "
58         << kernel->stats->totalTicks);
59 >     switch (which) {
264         ASSERTNOTREACHED();
265     }

```

Figure 5: ExceptionHandler()_SC_Halt_1

```

59     switch (which) {
60         case SyscallException:
61             switch (type) {
62                 case SC_Halt:
63                     DEBUG(dbgSys, "Shutdown, initiated by user program.\n");
64                     SysHalt();
65                     cout << "in exception\n";
66                     ASSERTNOTREACHED();
67                     break;

```

Figure 6: ExceptionHandler()_SC_Halt_2

```

16
17 void SysHalt() { kernel->interrupt->Halt(); }
18

```

Figure 7: SysHalt()

1.1.5 SysHalt()

At this point, the program enters the **kernel system call** layer to invoke a system call. (Please refer to Figure 7 for more details.)

1.1.6 Interrupt::Halt()

```
228 void Interrupt::Halt() {
229     #ifndef NO_HALT_STAT
230         cout << "Machine halting!\n\n";
231         cout << "This is halt\n";
232         kernel->stats->Print();
233     #endif
234     delete kernel; // Never returns.
235 }
```

Figure 8: Halt()

Finally, as shown in Figure 8, it prints detailed information about this system execution and deletes the kernel object, which signifies the shutdown of NachOS.

1.2 SC_Create

SC_Create is a kind of system call designed to create a file.

1.2.1 ExceptionHandler()

Since `Create()` is a type of **exception**, it should be handled in `ExceptionHandler()`. In this case, the program enters the **SC_Create** case on line 98, as shown in Figure 9.

Next, the program fetches the arguments provided by the user program by **reading register r4** and stores them in the variable `val`. Since `val` contains the memory address of the filename, the program **retrieves the actual filename address** and stores it in the variable `filename`, which is then passed to the kernel system call `SysCreate()`.

When the program returns from `SysCreate()`, it **writes the return value into register r2** and **increments the program counter** before returning.


```

98     case SC_Create:
99         val = kernel->machine->ReadRegister(4);
100         {
101             char *filename = &(kernel->machine->mainMemory[val]);
102             // cout << filename << endl;
103             status = SysCreate(filename);
104             kernel->machine->WriteRegister(2, (int)status);
105         }
106         kernel->machine->WriteRegister(
107             PrevPCReg, kernel->machine->ReadRegister(PCReg));
108         kernel->machine->WriteRegister(
109             PCReg, kernel->machine->ReadRegister(PCReg) + 4);
110         kernel->machine->WriteRegister(
111             NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
112         return;
113         ASSERTNOTREACHED();
114         break;

```

Figure 9: ExceptionHandler_SC_Create

```

31 int SysCreate(char *filename) {
32     // return value
33     // 1: success
34     // 0: failed
35     return kernel->fileSystem->Create(filename);
36 }

```

Figure 10: SysCreate()

1.2.2 SysCreate()

After entering the kernel system call layer, the program invokes the `Create()` system call in the `FileSystem` class, passing the filename to it. (Please refer to Figure 10 for more details.)

1.2.3 FileSystem::Create()

```
56     bool Create(char *name) {
57         int fileDescriptor = OpenForWrite(name);
58
59         if (fileDescriptor == -1) return FALSE;
60         Close(fileDescriptor);
61         return TRUE;
62     }
```

Figure 11: `FileSystem::Create()`

In this assignment, we use the **stuf** file system, which means we are actually calling the UNIX file system calls. As shown in Figure 11, the program calls `OpenForWrite()`, which internally invokes the UNIX `open()` function to obtain a file descriptor. The result, indicating success or failure, is then returned to the upper layer.

1.3 SC_PrintInt

The `SC_PrintInt` system call is designed to print integers to console display devices.

1.3.1 ExceptionHandler()

`ExceptionHandler()` is invoked when a user program makes a system call or raises an exception due to an illegal operation. This function serves as an intermediary, allowing the operating system to execute kernel-level code in response to events or requests from user programs.

System call arguments are passed from the user program to the kernel via specific CPU registers. The first step in handling the exception is to read register `r2` to determine the type of the system call and store into `type`, as shown on line 53. In this example, the value of `r2` is set to the constant `SC_PrintInt`, which is defined as 16 for the `PrintInt` system call. The function then checks if the exception type is a system call by examining the value of the `which` variable (line

```

50 void ExceptionHandler(ExceptionType which) {
51     char ch;
52     int val;
53     int type = kernel->machine->ReadRegister(2);
54     int status, exit, threadID, programID, fileID, numChar;
55     DEBUG(dbgSys, "Received Exception " << which << " type: " << type << "\n");
56     DEBUG(dbgTraCode, "In ExceptionHandler(), Received Exception "
57         << which << " type: " << type << ", "
58         << kernel->stats->totalTicks);
59     switch (which) {
60     case SyscallException:
61         switch (type) {
62         case SC_Halt:
63             DEBUG(dbgSys, "Shutdown, initiated by user program.\n");
64             SysHalt();
65             cout << "in exception\n";
66             ASSERTNOTREACHED();
67             break;
68         case SC_PrintInt:
69             DEBUG(dbgSys, "Print Int\n");
70             val = kernel->machine->ReadRegister(4);
71             DEBUG(dbgTraCode,
72                 "In ExceptionHandler(), into SysPrintInt, "
73                 << kernel->stats->totalTicks);
74             SysPrintInt(val);
75             DEBUG(dbgTraCode,
76                 "In ExceptionHandler(), return from SysPrintInt, "
77                 << kernel->stats->totalTicks);
78             // Set Program Counter
79             kernel->machine->WriteRegister(
80                 PrevPCReg, kernel->machine->ReadRegister(PCReg));
81             kernel->machine->WriteRegister(
82                 PCReg, kernel->machine->ReadRegister(PCReg) + 4);
83             kernel->machine->WriteRegister(
84                 NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
85             return;
86             ASSERTNOTREACHED();
87             break;

```

Figure 12: ExceptionHandler() in userprog/exception.cc

60). If which matches `SyscallException`, the function proceeds to handle system calls by entering a nested switch case based on the specific system call type stored in `type`.

In the case of `SC_PrintInt`, the argument for the this system call, the integers to be printed, is passed through register `r4` and stored in `val` (line 70). This value, `val`, is then passed as an argument to the function `SysPrintInt`, which executes the print operation.

The `dbgSys` of debugging messages is generated before obtaining the value of `val`, showing that the system is going to handle the print integer system call. As for the `dbgTraCode` of debugging messages, they are before and after the `SysPutInt()` call. These messages help track the execution flow of the function.

After executing `SysPrintInt`, the `ExceptionHandler()` function adjusts the program counter registers to resume execution of the user program on line 79 to 84. Specifically, it sets `PrevPCReg` to the current program counter (`PCReg`), then increments `PCReg` by 4 to point to the next instruction, and finally, it updates `NextPCReg` to the new value of `PCReg + 4`. This ensures that the user program continues executing from the instruction following the system call. If no matching case is found for the type of system call or exception, the function triggers an assertion failure, terminating execution due to an unhandled case (You may refer to Figure 12).

1.3.2 SysPrint()

```
19 void SysPrintInt(int val) {
20     DEBUG(dbgTraCode,
21         "In ksyscall.h:SysPrintInt, into synchConsoleOut->PutInt, "
22         "<< kernel->stats->totalTicks);
23     kernel->synchConsoleOut->PutInt(val);
24     DEBUG(dbgTraCode,
25         "In ksyscall.h:SysPrintInt, return from synchConsoleOut->PutInt, "
26         "<< kernel->stats->totalTicks);
27 }
```

Figure 13: `SysPrint()` in `userprog/ksyscall.h`

The `SysPrintInt()` function, as illustrated in Figure 13, provides a kernel-level interface for handling the `SC_PrintInt` system call, enabling user programs to print integer values to the console display device.

Within this function, the `PutInt()` method of the `synchConsoleOut` object is invoked, with the integer value `val` passed as an argument.

Additionally, debug messages are generated immediately before and after the `PutInt()` call, logging the current tick count using `kernel->stats->totalTicks`.

1.3.3 SynchConsoleOutput::PutInt()

```
100 void SynchConsoleOutput::PutInt(int value) {
101     char str[15];
102     int idx = 0;
103     // sprintf(str, "%d\n", value); the true one
104     sprintf(str, "%d\n\0", value); // simply for trace code
105     lock->Acquire();
106     do {
107         DEBUG(dbgTraCode, "In SynchConsoleOutput::PutInt, into consoleOutput->PutChar, " << kernel->stats->totalTicks);
108         consoleOutput->PutChar(str[idx]);
109         DEBUG(dbgTraCode, "In SynchConsoleOutput::PutInt, return from consoleOutput->PutChar, " << kernel->stats->totalTicks);
110         idx++;
111     } while (str[idx] != '\0');
112     while (str[idx] != '\0') {
113         DEBUG(dbgTraCode, "In SynchConsoleOutput::PutInt, into waitFor->P(), " << kernel->stats->totalTicks);
114         waitFor->P();
115         DEBUG(dbgTraCode, "In SynchConsoleOutput::PutInt, return from waitFor->P(), " << kernel->stats->totalTicks);
116     } while (str[idx] != '\0');
117     lock->Release();
118 }
```

Figure 14: SynchConsoleOutput::PutInt() in userprog/synchconsole.cc

This `SynchConsoleOutput::PutInt()` method, part of the `SynchConsoleOutput` class, manages the actual process of outputting the integers to the console in a synchronized manner, ensuring thread-safe and continuous output. You may refer to Figure 14.

The integer `value` is first converted into a character array `str` using `sprintf()`, with a newline character (`\n`) added for readability and a null terminator (`\0`) to mark the end of the string. Besides, the index `idx` keeps track of the current character being output.

A lock is acquired at the start of the function to prevent concurrent modifications during the output operation on line 105. In the while loop, characters are printed one by one by calling `PutChar()` on `ConsoleOutput`, with `idx` incremented after each call. After each character is printed, `waitFor->P()`, the consumer, is invoked to synchronize output with the availability of resources, likely ensuring the console is ready for the next character.

The loop continues until the null terminator (`\0`) is reached, at which point lock is released, allowing other processes to access.

1.3.4 ConsoleOutput::PutChar()

This `ConsoleOutput::PutChar()` method, part of the `ConsoleOutput` class, handles outputs a single character to the simulated console. You may refer to Figure 15.

It first verifies that no other character is currently being output by asserting `putBusy == FALSE`. If no other output is in progress, it writes the character `ch` to a file associated with the console using `WriteFile()`. The method then sets `putBusy` to `TRUE` to indicate that an output operation is active. Finally, it schedules an interrupt using `Schedule()`, specifying a delay of `ConsoleTime` and an

```

154 void ConsoleOutput::PutChar(char ch) {
155     ASSERT(putBusy == FALSE);
156     WriteFile(writeFileNo, &ch, sizeof(char));
157     putBusy = TRUE;
158     kernel->interrupt->Schedule(this, ConsoleTime, ConsoleWriteInt);
159 }

```

Figure 15: ConsoleOutput::PutChar() in machine/console.cc

interrupt type of ConsoleWriteInt to manage the timing of the next output (because of console's asynchronous nature). This delay and interrupt emulate the behavior of a hardware console that signals once output is complete (this is why the ConsoleOut class inherits from CallbackObj object-the ConsoleOut class allows the interrupt to call back once the output is finished).

1.3.5 Interrupt::Schedule()

```

289 void Interrupt::Schedule(CallbackObj *toCall, int fromNow, IntType type) {
290     int when = kernel->stats->totalTicks + fromNow;
291     PendingInterrupt *toOccur = new PendingInterrupt(toCall, when, type);
292
293     DEBUG(dbgInt, "Scheduling interrupt handler the " << intTypeNames[type] << " at time = " << when);
294     ASSERT(fromNow > 0);
295
296     pending->Insert(toOccur);
297 }

```

Figure 16: Interrupt::Schedule() in machine/interrupt.cc

The Interrupt::Schedule() simulates scheduling a hardware interrupt to occur at a specified future time. You may refer to Figure 16.

First, it calculates when, the simulated time at which the interrupt should occur, by adding fromNow to the current tick count (kernel->stats->totalTicks). Then, it creates a PendingInterrupt object, toOccur containing the callback object toCall, the timing when, and the interrupt type type. After verifying the delay fromNow is valid, it insert toOccur into pending interrupt list, where the interrupt will be processed when its scheduled time arrives.

1.3.6 Interrupt::CheckIfDue()

The Interrupt::CheckIfDue() method is called by Interrupt::Idle() to check if any scheduled interrupts are ready to execute. You may refer to Figure 17. When ConsoleOutput::PutChar() finishes, waitfor->P() triggers Thread::Sleep(), which then calls Interrupt::Idle(). Interrupt::Idle() verifies pending interrupts by calling Interrupt::CheckIfDue() with advanceClock

```

311 bool Interrupt::CheckIfDue(bool advanceClock) {
312     PendingInterrupt *next;
313     Statistics *stats = kernel->stats;
314
315     ASSERT(level == IntOff); // interrupts need to be disabled,
316                             // to invoke an interrupt handler
317     if (debug->IsEnabled(dbgInt)) {
318         DumpState();
319     }
320     if (pending->IsEmpty()) { // no pending interrupts
321         return FALSE;
322     }
323     next = pending->Front();
324
325     if (next->when > stats->totalTicks) {
326         if (!advanceClock) { // not time yet
327             return FALSE;
328         } else { // advance the clock to next interrupt
329             stats->idleTicks += (next->when - stats->totalTicks);
330             stats->totalTicks = next->when;
331             // UDelay(1000L); // rcgood - to stop nachos from spinning.
332         }
333     }
334
335     DEBUG(dbgInt, "Invoking interrupt handler for the ");
336     DEBUG(dbgInt, intTypeNames[next->type] << " at time " << next->when);
337
338     if (kernel->machine != NULL) {
339         kernel->machine->DelayedLoad(0, 0);
340     }
341
342     inHandler = TRUE;
343     do {
344         next = pending->RemoveFront(); // pull interrupt off list
345         DEBUG(dbgTraCode, "In Interrupt::CheckIfDue, into callOnInterrupt->CallBack, " << stats->totalTicks);
346         next->callOnInterrupt->CallBack(); // call the interrupt handler
347         DEBUG(dbgTraCode, "In Interrupt::CheckIfDue, return from callOnInterrupt->CallBack, " << stats->totalTicks);
348         delete next;
349     } while (!pending->IsEmpty() && (pending->Front()->when <= stats->totalTicks));
350     inHandler = FALSE;
351     return TRUE;
352 }

```

Figure 17: Interrupt::CheckIfDue() in machine/interrupt.cc

```

In SynchConsoleOutput::PutInt, into consoleOutput->PutChar, 66
1In SynchConsoleOutput::PutInt, return from consoleOutput->PutChar, 66
In SynchConsoleOutput::PutInt, into waitFor->P(), 66
In Semaphore::P(), 66      Invoked by SynchConsoleOutput::PutInt() calling waitFor->P()
In Thread::Sleep, Sleeping thread: con Invoked by Semaphore::P() calling currentThread->Sleep(FALSE)
In Interrupt::Idle, into CheckIfDue, 66 Invoked by Thread::Sleep() calling kernel->interrupt->Idle()
In Interrupt::CheckIfDue, into callOnInterrupt->C Invoked by Interrupt::Idle() calling CheckIfDue(TRUE)
In Interrupt::CheckIfDue, return from callOnInterrupt->CallBack, 100
In Interrupt::CheckIfDue, into callOnInterrupt->CallBack, 100
In Interrupt::CheckIfDue, return from callOnInterrupt->CallBack, 100
In Interrupt::CheckIfDue, into callOnInterrupt->CallBack, 100
In Interrupt::CheckIfDue, return from callOnInterrupt->CallBack, 100
In Interrupt::Idle, return true from CheckIfDue, 100 Back to Interrupt::Idle(), then back to Thread::Sleep()
In Interrupt::Idle, into CheckIfDue, 100 Invoked by Thread::Sleep() calling kernel->interrupt->Idle()
In Interrupt::CheckIfDue, into callOnInterrupt->C Invoked by Interrupt::Idle() calling CheckIfDue(TRUE)
In ConsoleOutput::CallBack(), 166
In SynchConsoleOutput::CallBack(), 166
In Semaphore::V(), 166
In Interrupt::CheckIfDue, return from callOnInterrupt->CallBack, 166
In Interrupt::Idle, return true from CheckIfDue, 166
In SynchConsoleOutput::PutInt, return from waitFor->P(), 176

```

Figure 18: Flow of calling Interrupt::CheckIfDue()

being TRUE (meaning the ready queue is empty). You may refer to Figure 18. These operations simulate a process requesting resources and periodically checking for interrupt callbacks.

First, `Interrupt::CheckIfDue()` ensures that interrupts are disabled to prevent concurrent executions. If debugging is enabled, it displays the current tick count and pending interrupts. It then checks if there are any pending interrupts. If not, it returns FALSE. In this example, pending interrupts exist, so it retrieves the next scheduled interrupt and checks if it is due by comparing `next->when` with `stats->totalTicks` to see if it is due. If `next->when` is greater than the current time and `advanceClock` is TRUE, it updates idle time and advances the clock to `next->when`.

The method then clears the CPU state for interrupt handling and sets TRUE to `inHandler`. It processes to remove each due interrupt from the list and invoke its interrupt handler by calling `next->callOnInterrupt->CallBack()`. This call triggers the specific callback associated with the interrupt, enabling the appropriate action to be executed-in this example, signaling that console output is complete. The loop continues until no more pending interrupts are due. Once all due interrupts are handled, `inHandler` is set back to FALSE to indicate that the system has exited interrupt handling, and the function returns TRUE.

1.3.7 ConsoleOutput::CallBack()

The `ConsoleOutput::CallBack()` method is triggered by `Interrupt::CheckIfDue()` to signal that `ConsoleOutput::PutChar()` has completed. You may refer to Fig-


```

141 void ConsoleOutput::CallBack() {
142     DEBUG(dbgTraCode, "In ConsoleOutput::CallBack(), " << kernel->stats->totalTicks);
143     putBusy = FALSE;
144     kernel->stats->numConsoleCharsWritten++;
145     callWhenDone->CallBack();
146 }

```

Figure 19: ConsoleOutput::CallBack() in machine/console.cc

ure 19.

The method first sets `putBusy` back to `FALSE`, indicating that the console is ready for the next character. It then increment `kernel->stats->numConsoleCharsWritten` to update the count of characters successfully output. Finally, it calls `callWhenDone->CallBack()`, which signals the next character can process.

1.3.8 SynchConsoleOutput::CallBack()

```

131 void SynchConsoleOutput::CallBack() {
132     DEBUG(dbgTraCode,
133         "In SynchConsoleOutput::CallBack(), " << kernel->stats->totalTicks);
134     waitFor->V();
135 }

```

Figure 20: SynchConsoleOutput::CallBack() in userprog/console.cc

The `SynchConsoleOutput::CallBack()` method is triggered by `Interrupt::CheckIfDue()` to signal it is safe to send the next character to the console. You may refer to Figure 20.

The method calls `waitFor->V()`, the producer, which releases the semaphore, increasing the availability of resources-in this case, allowing the next character to be output to the console.

1.3.9 Machine::Run()

The `Machine::Run()` method simulates the execution of a user-level program, processing instructions one by one. You may refer to Figure 21.

It first creates `instr` to hold each decoded instruction. After setting the execution mode to user mode, it enters an infinite for loop. In each loop iteration, `OneInstruction(instr);` is called to execute a single instruction, followed by `OneTick()` to advance one clock. If single-step debugging is enabled and the specified tick count is reached, `Debugger()` is called to display debug messages.

```

54 void Machine::Run() {
55     Instruction *instr = new Instruction; // storage for decoded instruction
56     if (debug->IsEnabled('m')) {
57         cout << "Starting program in thread: " << kernel->currentThread->getName();
58         cout << ", at time: " << kernel->stats->totalTicks << "\n";
59     }
60     kernel->interrupt->setStatus(UserMode);
61     for (;;) {
62         DEBUG(dbgTraCode, "In Machine::Run(), into OneInstruction "
63             << "==" Tick " << kernel->stats->totalTicks << " ==");
64         OneInstruction(instr);
65         DEBUG(dbgTraCode, "In Machine::Run(), return from OneInstruction "
66             << "==" Tick " << kernel->stats->totalTicks << " ==");
67
68         DEBUG(dbgTraCode, "In Machine::Run(), into OneTick "
69             << "==" Tick " << kernel->stats->totalTicks << " ==");
70         kernel->interrupt->OneTick();
71         DEBUG(dbgTraCode, "In Machine::Run(), return from OneTick "
72             << "==" Tick " << kernel->stats->totalTicks << " ==");
73         if (singleStep && (runUntilTime <= kernel->stats->totalTicks))
74             Debugger();
75     }
76 }

```

Figure 21: Machine::Run() in machine/mipssim.cc

1.3.10 Interrupt::OneTick()

The Interrupt::OneTick() method is called by Machine::Run() to advance the simulated clock by one tick. You may refer to Figure 22.

The method first updates stats->totalTicks and either stats->systemTicks or stats->userTicks, depending on the current execution mode. As time moves on, it is necessary to verify if any pending interrupts are due. The method then temporarily disables interrupts to check for any pending interrupts that are due by calling CheckIfDue(FALSE) (since it supposes there exist instructions not executed), re-enabling interrupts afterward.

The second thing to check is whether any interrupt handlers call back and request a context switching using yieldOnReturn. If yieldOnReturn is set, the method sets the mode to the kernel mode, SystemMode, performs the context switching and then restores the previous mode.

1.4 Makefile

A Makefile is a special file that contains paths to various resources, shell commands, and all the dependency records. You can think of it as a large shell script which can be executed partly by using make command. We explain the functionality of the Makefile by dividing it into three parts: variable definitions, shell com-

```

145 void Interrupt::OneTick() {
146     MachineStatus oldStatus = status;
147     Statistics *stats = kernel->stats;
148
149     // advance simulated time
150     if (status == SystemMode) {
151         stats->totalTicks += SystemTick;
152         stats->systemTicks += SystemTick;
153     } else {
154         stats->totalTicks += UserTick;
155         stats->userTicks += UserTick;
156     }
157     DEBUG(dbgInt, "== Tick " << stats->totalTicks << " ==");
158
159     // check any pending interrupts are now ready to fire
160     ChangeLevel(IntOn, IntOff); // first, turn off interrupts
161                                // (interrupt handlers run with
162                                // interrupts disabled)
163     CheckIfDue(FALSE);         // check for pending interrupts
164     ChangeLevel(IntOff, IntOn); // re-enable interrupts
165     if (yieldOnReturn) {        // if the timer device handler asked
166                                // for a context switch, ok to do it now
167         yieldOnReturn = FALSE;
168         status = SystemMode; // yield is a kernel routine
169         kernel->currentThread->Yield();
170         status = oldStatus;
171     }
172 }

```

Figure 22: Interrupt::OneTick() in machine/interrupt.cc

mands and dependencies, and the clean-up and error messages.

1.4.1 Variable Definitions

```
102 include Makefile.dep
103
104 CC = $(GCCDIR)gcc
105 AS = $(GCCDIR)as
106 LD = $(GCCDIR)ld
107
108 INCDIR = ../userprog -I../lib
109 CFLAGS = -g -O 0 -c $(INCDIR) -B/usr/bin/local/nachos/lib/gcc-lib/decstation-ultrix/2.95.2/ -B/usr/bin/local/nachos/decstation-ultr
110
111 ifeq ($(hosttype),unknown)
112 PROGRAMS = unknownhost
113 else
114 # change this if you create a new test program!
115 PROGRAMS = add halt createFile LotOfAdd
116 endif
117
```

Figure 23: MakeFile: Variable Definitions

As shown in Figure 23, at the beginning, the Makefile includes Makefile.dep, which defines many paths using variables. Next, from lines 104 to 106, it uses the variables defined in Makefile.dep to set the paths for several important tools, such as the compiler, assembler, and linker, in the new variables CC, AS, and LD, respectively. From lines 108 to 109, the Makefile defines INCDIR as the path where the compiler can find the included header files, and CFLAGS as part of the command that specifies how the compiler should compile the files.

1.4.2 Shell Commands and Dependencies

```
191 fileIO_test1.o: fileIO_test1.c
192 | $(CC) $(CFLAGS) -c fileIO_test1.c
193 fileIO_test1: fileIO_test1.o start.o
194 | $(LD) $(LDFLAGS) start.o fileIO_test1.o -o fileIO_test1.coff
195 | $(COFF2NOFF) fileIO_test1.coff fileIO_test1
196
197 fileIO_test2.o: fileIO_test2.c
198 | $(CC) $(CFLAGS) -c fileIO_test2.c
199 fileIO_test2: fileIO_test2.o start.o
200 | $(LD) $(LDFLAGS) start.o fileIO_test2.o -o fileIO_test2.coff
201 | $(COFF2NOFF) fileIO_test2.coff fileIO_test2
```

Figure 24: MakeFile: Shell Commands and Dependencies

Since there are many similar code files, we will use fileIO_test1.c and fileIO_test2.c as examples to explain as shown in Figure 24.

The process can be divided into two parts: compiling and linking. For the compiling part, from lines 191 to 192, the Makefile uses the variables defined earlier, CC and CFLAGS, to compile fileIO_test1.c into fileIO_test1.o. For the linking part, from lines 193 to 195, the Makefile uses several variables to link

fileIO_test1.o and start.o into an executable file, fileIO_test1. The process for fileIO_test2 is the same as for fileIO_test1.

1.4.3 Clean-Up and Error Messages

```
233 clean:
234     $(RM) -f *.o *.ii
235     $(RM) -f *.coff
236
237 distclean: clean
238     $(RM) -f $(PROGRAMS)
239
240 unknownhost:
241     @echo Host type could not be determined.
242     @echo make is terminating.
243     @echo If you are on an MFCF machine, contact the instructor to report this problem
244     @echo Otherwise, edit Makefile.dep and try again.
```

Figure 25: MakeFile: Clean-Up and Error Messages

clean and distclean are two targets defined in the Makefile. These targets are actually commands made up of several predefined variables. The clean target is used to remove all .o, .ii, and .coff files, while distclean is used to clean up the entire executable program. Lastly, unknownhost is a target used to output error messages.

2 System Call Implementation

In this section, we implement four I/O system calls for NachOS: SC_Open, SC_Read, SC_Write, and SC_Close. We begin by examining their usage in fileIO_test1 and fileIO_test2, referencing the #define definitions in userprog/syscall.h where these system calls are initially commented out.

Using these files and the hints from the course slide, we observe that the flow of these I/O system calls follows a similar pattern, as illustrated in Figure 26. First, we explain the common steps (the first two steps) in the flow for all four I/O system call, then we detail the unique aspects of each call separately.

First of all, we uncomment lines 27, 28, 29, and 31 to enable the system calls SC_Open, SC_Read, SC_Write, and SC_Close in userprog/syscall.h. You may refer to Figure 27.

Subsequently, we add assembly code for the Open, Write, Read, and Close system calls, following the existing system calls' logic. Each system calls first loads its corresponding system call value into register r2, executes the system call using the syscall instruction, and finally returns to the caller's address. You may refer to Figure 28.

userprog/syscall.h	<pre> #define SC_Open 6 #define SC_Read 7 #define SC_Write 8 #define SC_Close 10 </pre>
test/start.S	<pre> /* add Open, Write, Read, and Close assembly code */ </pre>
userprog/exception.cc	<pre> ExceptionHandler() </pre>
userprog/ksyscall.h	<pre> SysOpen() SysWrite() SysRaed() SysClose() </pre>
filesys/filesys.h	<pre> OpenAFile() WriteFile() ReadFile() CloseFile() </pre>

Figure 26: Flow of implementing the four I/O system calls

```

21  #define SC_Halt 0
22  #define SC_Exit 1
23  #define SC_Exec 2
24  #define SC_Join 3
25  #define SC_Create 4
26  #define SC_Remove 5
27  #define SC_Open 6
28  #define SC_Read 7
29  #define SC_Write 8
30  #define SC_Seek 9
31  #define SC_Close 10
32  #define SC_ThreadFork 11
33  #define SC_ThreadYield 12
34  #define SC_ExecV 13
35  #define SC_ThreadExit 14
36  #define SC_ThreadJoin 15
37  #define SC_PrintInt 16
38  #define SC_Add 42
39  #define SC_MSG 100
40  #ifndef IN_ASM

```

Figure 27: Removing the comments from the code

```

46  Open:
47      addiu $2,$0,SC_Open
48      syscall
49      j    $31
50      .end Open
51
52      .globl Write
53      .ent    Write
54  Write:
55      addiu $2,$0,SC_Write
56      syscall
57      j    $31
58      .end Write
59
60      .globl Read
61      .ent    Read
62  Read:
63      addiu $2,$0,SC_Read
64      syscall
65      j    $31
66      .end Read
67
68      .globl Close
69      .ent    Close
70  Close:
71      addiu $2,$0,SC_Close
72      syscall
73      j    $31
74      .end Close
75
76      .globl Halt
77      .ent    Halt

```

Figure 28: Adding assembly code in Start.s

2.1 Open

2.1.1 ExceptionHandler()

```
116 case SC_Open:
117     /*
118     #define SC_Open 6
119     Open a file for read & write.
120     */
121     DEBUG(dbgTrCode, "In ExceptionHandler:case SC_Open.");
122     val = kernel->machine->ReadRegister(4); // Retrieve file name address.
123     {
124         char *filename = &(kernel->machine->mainMemory[val]); // Retrive file name.
125         DEBUG(dbgTrCode, "In ExceptionHandler:case SC_Open, into SysOpen.");
126         fileID = SysOpen(filename); // Success: get file ID / Fail: get -1
127         DEBUG(dbgTrCode, "In ExceptionHandler:case SC_Open, return from SysOpen.");
128         kernel->machine->WriteRegister(2, fileID); // Write file ID into register.
129     }
130     kernel->machine->WriteRegister(
131         | PrevPCReg, kernel->machine->ReadRegister(PCReg));
132     kernel->machine->WriteRegister(
133         | PCReg, kernel->machine->ReadRegister(PCReg) + 4);
134     kernel->machine->WriteRegister(
135         | NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
136     return;
137     ASSERTNOTREACHED();
138     break;
139
```

Figure 29: ExceptionHandler Case: SC_Open

In `ExceptionHandler()`, we add a case for `SC_Open`. You may refer to Figure 29.

First, we retrieve the system call argument stored in register `r4`, which holds the address of the file name in memory, on line 122. To avoid duplicate assignment, we use a block to define `filename` and `fileID`. The file name the user want to open is obtained by accessing the physical memory at the address from `r4`. The file name is then passed to the `SysOpen()`, which returns an opened file ID (`fileID`). We store this opened file id in register `r2`, as `r2` is the return value register in MIPS.

Finally, we follows the existing code to advance program counter's value and perform necessary checks.

2.1.2 SysOpen()

```
39 OpenFileId SysOpen(char *name) {
40     | return kernel->fileSystem->OpenAFile(name);
41 }
42
```

Figure 30: `ksyscall.h/SysOpen()`

We uncomment the `SysOpen()` function, which calls `OpenAFile()` in `FileSystem`, to handle the actual Open system call. You may refer to Figure 30.

2.1.3 OpenFile()

```
73  OpenFileId OpenFile(char *name) {
74      DEBUG(dbgTraCode, "In FileSystem::OpenFile(), file name = " << name);
75      // Handling duplicate file opening
76      for (int idx = 0; idx < 20; ++idx) {
77          if ((OpenFileNames[idx] != NULL) &&
78              !strcmp(OpenFileNames[idx], name)) {
79              DEBUG(dbgTraCode,
80                  | "In FileSystem::OpenFile(), duplicate opening is found: " << OpenFileNames[idx]);
81              return -1;
82          }
83      }
84      // Searching an empty spot to place the opened file control object ptr &
85      // to record file name
86      for (int idx = 0; idx < 20; ++idx) {
87          if (OpenFileTable[idx] == NULL) {
88              // Get file fd
89              int fileDescriptor = OpenForReadWrite(name, FALSE);
90              DEBUG(dbgTraCode, "In FileSystem::OpenFile(), file descriptor(fd): " << fileDescriptor);
91              // Handling non-existent file
92              if (fileDescriptor == -1) return -1;
93
94              OpenFileTable[idx] = new OpenFile(fileDescriptor);
95              DEBUG(dbgTraCode,
96                  | "Created a new OpenFile and stored into OpenFileTable");
97              OpenFileNames[idx] = name;
98              DEBUG(dbgTraCode,
99                  | "Created a new OpenFile and stored its name into "
100                  | "OpenFileNames");
101              return idx;
102          }
103      }
104      // Handling exceeded file opening (at most 20 files)
105      return -1;
106  }
```

Figure 31: filesystem.h/OpenFile()

Before implementing `OpenFile()`, we define a new array, `OpenFileNames`, to store the names of opened files. The `OpenFile()` method returns `-1` to indicate a failure to open the file; otherwise, it returns the file ID, which is the index of the file in `OpenFileTable`. You may refer to Figure 31.

We first verify if the file name has already been opened. If a duplicate is found, it returns `-1`. If not, we call a for loop to find an empty spot to locate an `OpenFile` object pointer. If there is an empty space, we call `OpenForReadWrite` with the file name to obtain a file descriptor (the `FALSE` passed is imitating it in `Open()` method). We then check if the file descriptor is valid. There are two cases: the file does not exist or the maximum limit of open files (20) is exceeded. If the file does exist, we store a new `OpenFile` object and record the file name at the same index in `OpenFileNames` then return the index (`idx`) as the opened file ID.

Figure 31 shows the complete code of our implementation.

```

140 case SC_Write:
141     /*
142     #define SC_Write 8
143     Write "size" characters from the buffer into the file, and
144     return the number of characters actually written to the file
145     */
146     DEBUG(dbgTrcCode, "In ExceptionHandler:case SC_Write.");
147     val = kernel->machine->ReadRegister(4);
148     {
149         char *buffer = &(kernel->machine->mainMemory[val]);
150         numChar = kernel->machine->ReadRegister(5);
151         fileID = kernel->machine->ReadRegister(6);
152
153         DEBUG(dbgTrcCode, "In ExceptionHandler:case SC_Write, into SysWrite()");
154         numChar = SysWrite(buffer, numChar, fileID); // Success: get number of character written into the file / Fail: get -1
155         DEBUG(dbgTrcCode, "In ExceptionHandler:case SC_Write, return from SysWrite()");
156         kernel->machine->WriteRegister(2, numChar); // Write result into register.
157     }
158     kernel->machine->WriteRegister(
159         PrevPCReg, kernel->machine->ReadRegister(PCReg));
160     kernel->machine->WriteRegister(
161         PCReg, kernel->machine->ReadRegister(PCReg) + 4);
162     kernel->machine->WriteRegister(
163         NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
164     return;
165     ASSERTNOTREACHED();
166     break;
167

```

Figure 32: ExceptionHandler Case: SC_Write

2.2 Write

2.2.1 ExceptionHandler()

In `ExceptionHandler()`, we add a case for `SC_Open`. You may refer to Figure 32.

Similar to the `SC_Open` case, `SC_Write` case handles additional arguments, the number of characters to write (`numChar`) and the opened file ID (`fileID`), stored in register `r5` and `r6`, respectively. Instead of a file name, we access the physical memory at the address in `r4` to retrieve the buffer containing the data to write. The `numChar` and `fileID` arguments are stored directly in registers probably since they are small enough to fit.

After calling `SysWrite()`, we store its return value, the number of characters actually written, into register `r2`.

2.2.2 SysWrite()

```

43 int SysWrite(char *name, int size, OpenFileId id) {
44     | return kernel->fileSystem->WriteFile(name, size, id);
45 }
46

```

Figure 33: `ksyscall.h/SysWrite()`

In `ksyscall.h`, we add a function `SysWrite()`. In `SysWrite()`, we call `WriteFile()`, which is also in `FileSystem`, to handle the actual Write system call. You may refer to Figure 33.

2.2.3 WriteFile()

```
110 int WriteFile(char *buffer, int size, OpenFileId id) {
111     DEBUG(DBG_TraCode, "In FileSystem::WriteFile()");
112
113     if (id < 0 || id >= 20) { // Handling invalid file ID
114         return -1;
115     } else {
116         if (OpenFileTable[id] == NULL) { // Handling file not yet opened
117             return -1;
118         }
119     }
120     return OpenFileTable[id]->Write(buffer, size * sizeof(char));
121 }
```

Figure 34: filesystem/WriteFile()

In `filesystem.h`, we uncomment the `WriteFile()` method. This method returns `-1` to indicate a failure to write the file; otherwise, it returns the number of characters written. You may refer to Figure 34.

In `WriteFile()`, we first check the file ID is within the valid range (0 to 19). If the ID is valid, we then verify the file is open by checking `OpenFileTable[id]`. If the file is open, we call `Write` method, passing `buffer` and `size * sizeof(char)` (size of the characters in memory) as arguments.

You may refer to 34 for details.

2.3 Read

2.3.1 ExceptionHandler()

In `ExceptionHandler()`, we add a case for `SC_Read`. You may refer to Figure 35.

Similar to the `SC_Write` case, the `SC_Read` case calls `SysRead()` instead. We store its return value, the number of characters actually read, into register `r2`.

2.3.2 SysRead()

In `ksyscall.h`, we add a function `SysRead()`. In `SysRead()`, we call `ReadFile()`, which is also in `FileSystem`, to handle the actual `Read` system call. You may refer to Figure 36.

2.3.3 ReadFile()

In `filesystem.h`, we uncomment the `ReadFile()` method. You may refer to Figure 37.

```

168 case SC_Read:
169     /*
170     #define SC_Read 7
171     Read "size" characters from the file into the buffer, and
172     return the number of characters actually read from the file
173     */
174     DEBUG(dbgTraCode, "In ExceptionHandler:case SC_Read.");
175     val = kernel->machine->ReadRegister(4);
176     {
177         char *buffer = &(kernel->machine->mainMemory[val]);
178         numChar = kernel->machine->ReadRegister(5);
179         fileId = kernel->machine->ReadRegister(6);
180
181         DEBUG(dbgTraCode, "In ExceptionHandler:case SC_Read, into SysRaed()");
182         numChar = SysRaed(buffer, numChar, fileId);
183         DEBUG(dbgTraCode, "In ExceptionHandler:case SC_Read, return from SysRaed()");
184         kernel->machine->WriteRegister(2, numChar);
185     }
186     kernel->machine->WriteRegister(
187         PrevPCReg, kernel->machine->ReadRegister(PCReg));
188     kernel->machine->WriteRegister(
189         PCReg, kernel->machine->ReadRegister(PCReg) + 4);
190     kernel->machine->WriteRegister(
191         NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
192     return;
193     ASSERTNOTREACHED();
194     break;
195

```

Figure 35: ExceptionHandler Case: SC_Read

```

47 int SysRaed(char *name, int size, OpenFileId id) {
48     Click to add a breakpoint kernel->fileSystem->ReadFile(name, size, id);
49 }
50

```

Figure 36: ksyscall.h/SysRead()

```

123 int ReadFile(char *buffer, int size, OpenFileId id) {
124     DEBUG(dbgTraCode, "In FileSystem::RaedFile()");
125     if (id < 0 || id >= 20) { // Handling invalid file ID
126         return -1;
127     } else {
128         if (OpenFileTable[id] == NULL) { // Handling file not yet opened
129             return -1;
130         }
131     }
132     return OpenFileTable[id]->Read(buffer, size * sizeof(char));
133 }

```

Figure 37: filesys.h/ReadFile()

Similar to `WriteFile()`, this method verifies the file ID and checks the file is open. If these checks pass, it calls the `Read` method on the file instead, passing `buffer` and `size * sizeof(char)` as arguments.

You may refer to 37 for details.

2.4 Close

2.4.1 ExceptionHandler()

```

196 | case SC_Close:
197 |     /*
198 |     #define SC_Close 10
199 |     Close the file.
200 |     */
201 |     DEBUG(dbgTracode, "In ExceptionHandler:case SC_Close.");
202 |     val = kernel->machine->ReadRegister(4);
203 |     DEBUG(dbgTracode, "In ExceptionHandler:case SC_Close, into SysClose()");
204 |     status = SysClose(val);
205 |     DEBUG(dbgTracode, "In ExceptionHandler:case SC_Close, return from SysClose()");
206 |     kernel->machine->WriteRegister(2, status);
207 |     kernel->machine->WriteRegister(
208 |         | PrevPCReg, kernel->machine->ReadRegister(PCReg));
209 |     kernel->machine->WriteRegister(
210 |         | PCReg, kernel->machine->ReadRegister(PCReg) + 4);
211 |     kernel->machine->WriteRegister(
212 |         | NextPCReg, kernel->machine->ReadRegister(PCReg) + 4);
213 |     return;
214 |     ASSERTNOTREACHED();
215 |     break;

```

Figure 38: ExceptionHandler Case: SC_Close

In `ExceptionHandler()`, we add a case for `SC_Close`. You may refer to Figure 38.

Similar to the `SC_Open` case, the `SC_Close` case retrieves the value stored in register `r4` directly as the opened file ID, `val`. This ID is then passed to `SysClose()`, which returns a `status` indicating whether the file was successfully closed (1 for success, -1 for failure).

2.4.2 SysClose()

```

51 | int SysClose(OpenFileId id) {
52 |     | return kernel->fileSystem->CloseFile(id);
53 | }

```

Figure 39: `ksyscall.h/SysClose()`

In `ksyscall.h`, we add a function `SysClose()`. In `SysClose()`, we call `CloseFile()`, which also in `FileSystem`, to handle the actual `Close` system call. You may refer to Figure 39.

```

135     int CloseFile(OpenFileId id) {
136         DEBUG(dbgTraCode, "In FileSystem::CloseFile()");
137         if (id < 0 || id > 20) { // Handling invalid file ID
138             return -1;
139         } else {
140             if (OpenFileTable[id] == NULL) { // Handling file not yet opened
141                 return -1;
142             }
143         }
144         delete OpenFileTable[id];
145         OpenFileTable[id] = NULL;
146         OpenFileNames[id] = NULL;
147         return 1;
148     }

```

Figure 40: filesystem/CloseFile()

2.4.3 CloseFile()

In filesystem.h, we uncomment the CloseFile() method. You may refer to Figure 40.

Similar to WriteFile(), this method first verifies the file ID and checks the file is open. If these checks pass, we use delete to call its Close method and close the file. We then set OpenFileTable[id] and OpenFileNames[id] to NULL to free the slot for other files. You may refer to 40 for details.

To confirm the file is closed, we check if calling Length() on it after the deletion, which results in a segmentation fault, as illustrated in Figure 41a. In contrast, Figure 41b shows the program running successfully without calling Length() after deletion. Therefore, we probably could suggest that the file is indeed closed.

3 Difficulties

SC_PrintInt : In this process, we struggled to identify the timing of character output and how Interrupt::CheckIfDue() is invoked. Initially, we used add.c to print debug messages, but the small number of characters made it difficult to trace the flow. To address this, we created consoleIO_test3.c, as shown in Figure 42, to output more characters. This test helped us pinpoint when Interrupt::CheckIfDue() is called and understand the subsequent flow.

Open : When implementing the Open() system call, we would like to make sure we handled the file opening and closing process correctly by opening 20 files, then closing one file, and finally opening another file. While reviewing our debug messages, we discovered something tricky about the Unix Open() system call.

```

145     DEBUG(dbgTraCode,
146           "Before deletion, file length: " << OpenFileTable[id]->Length());
147     delete OpenFileTable[id];
148     OpenFileTable[id] = NULL;
149     OpenFileNames[id] = NULL;
150     DEBUG(dbgTraCode,
151           "After deletion, file length: " << OpenFileTable[id]->Length());
152     return 1;
153 }

```

PROBLEMS OUTPUT TERMINAL PORTS COMMENTS

TERMINAL

```

In Machine:Run(), return from OneInstruction == Tick 940 ==
In Machine:Run(), into OneTick == Tick 940 ==
In Machine:Run(), return from OneTick == Tick 941 ==
In Machine:Run(), into OneInstruction == Tick 941 ==
In Machine:OneInstruction, RaiseException(SyscallException, 0), 941
In ExceptionHandler(), Received Exception 1 type: 10, 941
In ExceptionHandler:case SC_close,
In ExceptionHandler:case SC_close, into SysClose()
In Filesystem:CloseFile()
Before deletion, file length: 26
./check.sh: line 10: 17338 Segmentation fault ./build.linux/nachos -d c -e fileI0_test1

```

(a) Call of Length() after the deletion

```

146     "Before deletion, file length: " << OpenFileTable[id]->Length());
147     delete OpenFileTable[id];
148     OpenFileTable[id] = NULL;
149     OpenFileNames[id] = NULL;
150     return 1;
151 }
152
153 bool Remove(char *name) { return Unlink(name) == 0; }
154

```

PROBLEMS OUTPUT TERMINAL PORTS COMMENTS

TERMINAL

```

In Machine:Run(), return from OneInstruction == Tick 952 ==
In Machine:Run(), into OneTick == Tick 952 ==
In Machine:Run(), return from OneTick == Tick 953 ==
In Machine:Run(), into OneInstruction == Tick 953 ==
In Machine:Run(), return from OneInstruction == Tick 953 ==
In Machine:Run(), into OneTick == Tick 953 ==
In Machine:Run(), return from OneTick == Tick 954 ==
In Machine:Run(), into OneInstruction == Tick 954 ==
In Machine:OneInstruction, RaiseException(SyscallException, 0), 954
In ExceptionHandler(), Received Exception 1 type: 100, 954
Success on creating fileI.test

```

(b) No call of Length() after the deletion

Figure 41: Result of two cases

```

NachOS-4.0_MP1 > code > test > C consoleIO_test3.c > main()
1  #include "syscall.h"
2
3  int main() {
4      int n = 1126;
5      PrintInt(n);
6
7      // return 0;
8      Halt();
9  }
10

```

Figure 42: Difficulties: SC_PrintInt

```

86     for (int idx = 0; idx < 20; ++idx) {
87         if (OpenFileTable[idx] == NULL) {
88             // Get file fd
89             int fileDescriptor = OpenForReadWrite(name, FALSE);
90             DEBUG(dbgTrCode, "In FileSystem::OpenFile(), file descriptor(fd):" << fileDescriptor);
91             // Handling non-existent file
92             if (fileDescriptor == -1) return -1;
93
94             OpenFileTable[idx] = new OpenFile(fileDescriptor);
95             DEBUG(dbgTrCode,
96                 | "Created a new OpenFile and stored into OpenFileTable");
97             OpenFileNames[idx] = name;
98             DEBUG(dbgTrCode,
99                 | "Created a new OpenFile and stored its name into "
100                | "OpenFileNames");
101             return idx;
102         }
103     }
104     // Handling exceeded file opening (at most 20 files)
105     return -1;
106 }

```

Figure 43: Difficulties: SC_Open Code

You may first refer to Figure 43 for the details of our code. Initially, we opened file1.test, and then we opened another 19 files, from test0.test to test18.test (these files have been created before we tried to open them). So far, the debug messages look great, every opened file had its own file descriptor as shown in Figure 44. Afterwards, we tried to close test18.test, and the debug message we added in OpenFile destructor (as shown in Figure 45) showed that we had successfully close the file (as shown in Figure 46). When we tried to opened test19.test (this file had also been created before), we found out that the returned value of `Open()` (i.e. file descriptor) is -1, you may refer to Figure 47. As shown in our implementation code (Figure 43), it means that the return value of test19.test was returned on line 92, this means that we indeed closed the file, controlled our `OpenFileTable` (because it went into line 87) but we failed to call `Open()` on our 21st try! It was supposed to be returned on line 107!

You may first refer to Figure 43 for the details of our code. Initially, we opened file1.test, followed by another 19 files, from test0.test to test18.test (these files had been created prior to opening them). Up to this point, the debug messages looked great—each file had its own file descriptor, as shown in Figure 44. Next, we closed test18.test, and the debug message from the OpenFile destructor (shown in Figure 45) confirmed that the file had been successfully closed, as shown in Figure 46.

However, when we tried to open test19.test (which had also been created before), we discovered that the returned value of `Open()` (i.e., the file descriptor) was -1. You may refer to Figure 47. According to our implementation code (Figure 43), this return value was produced on line 92, indicating that while we successfully closed the file and managed our `OpenFileTable` (as the code reached line 87), we failed to call `Open()` on our 21st attempt!

The file descriptor was expected to be returned on line 105.

```
In FileSystem::OpenAFile(), file name = test18.test  
In FileSystem::OpenAFile(), file descriptor(fd): 30
```

Figure 44: Difficulties: SC_Open Debug Message 1

```
37 ~OpenFile() {  
38     int retval = Close(file);  
39     DEBUG(dbgTraCode, "In OpenFile::destructor with close return value: " << retval);  
40     // close the file
```

Figure 45: Difficulties: SC_Open Debug Message 2

```
In FileSystem::CloseFile()  
In OpenFile::destructor with close return value: 0  
In ExceptionHandler:case SC_Close, return from SysClose()
```

Figure 46: Difficulties: SC_Open Debug Message 3

Write : We encountered some difficulties while implementing the `Write()` system call. Initially, we were puzzled as to why we couldn't fetch the arguments, `numChar` and `fileID`, by passing the address obtained from the register into memory. This issue troubled us for quite some time. You may refer to the code shown in Figure 48. Eventually, we resolved the problem by printing debug messages and discovered that the values of these arguments were actually stored in the registers. We didn't need to fetch them from memory; instead, we could load them directly from the registers and use them immediately.

After several days, we realized that we had forgotten to check the types of variables passed by the user in `fileIO_test1`. The arguments passed to `Write()` included an address (pointer) and two integers. Naturally, we could load `numChar` and `fileID` directly from the registers since they were not addresses!

Close : In the `CloseFile()` method, we want to confirm that the target `OpenFile` object is truly deleted. However, since the `OpenFile` destructor does not return a value, we cannot use `retval = delete OpenFileTable[id];` to verify deletion through a debug message. As a result, we decided to consult the TA to ensure our implementation is correct.

```
In FileSystem::OpenAFile(), file name = file19.test  
In FileSystem::OpenAFile(), file descriptor(fd): -1
```

Figure 47: Difficulties: SC_Open Debug Message 4

```
140 case SC_Write:  
141 /*  
142 #define SC_Write 8  
143 Write "size" characters from the buffer into the file, and  
144 return the number of characters actually written to the file  
145 */  
146 DEBUG(dbgTraCode, "In ExceptionHandler:case SC_Write.");  
147 val = kernel->machine->ReadRegister(4);  
148 {  
149     char *buffer = &(kernel->machine->mainMemory[val]);  
150     numChar = kernel->machine->ReadRegister(5);  
151     fileId = kernel->machine->ReadRegister(6);
```

Figure 48: Difficulties: SC_Write