

Operating System

MP2: Multi-Programming

Team Member(Team 10): 林子謙 (112065517)、陳奕潔
(110000212)

Contributions: Both contributed to page table implementing and report writing.

Contents

1	Code Tracing	2
1.1	Kernel::Kernel()	2
1.2	Initialization	2
1.2.1	Kernel::ExecAll()	2
1.2.2	Kernel::Exec()	5
1.2.3	AddrSpace::AddrSpace()	5
1.2.4	Thread::Fork()	5
1.2.5	Thread::StackAllocate()	6
1.2.6	Scheduler::ReadyToRun()	6
1.2.7	Thread::Finish()	6
1.2.8	Thread::Sleep()	6
1.2.9	Scheduler::Run()	6
1.3	Executing a Thread	10
1.3.1	Thread::Begin()	10
1.3.2	Scheduler::CheckToBeDestroyed()	10
1.3.3	ForkExecute()	10
1.3.4	AddrSpace::Load()	11
1.3.5	AddrSpace::Execute()	11
1.4	Questions in Spec	13
2	Page Table Implementation	14
2.1	Construct a page table	15
2.2	Assign the available frames	15
2.3	Load the code	19
2.4	Return the frames	21

1 Code Tracing

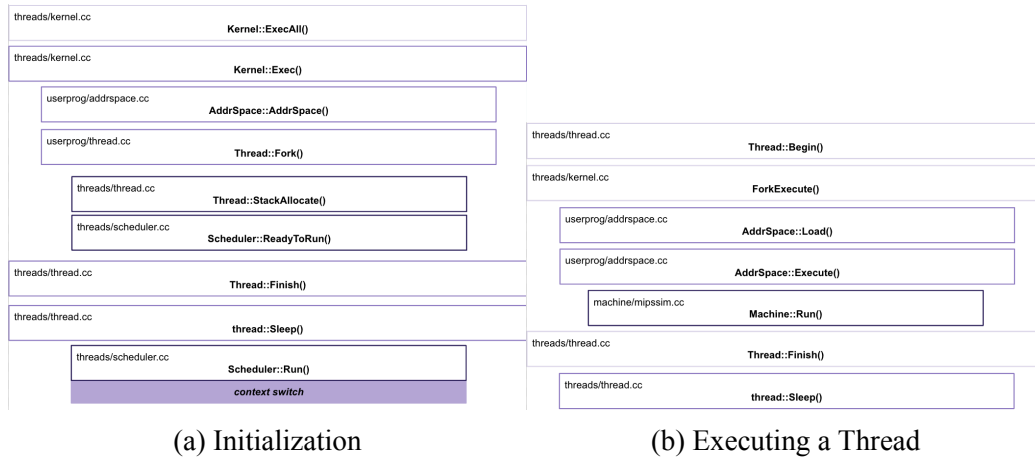


Figure 1: Code Flow

In this assignment, we begin tracing the code from `Kernel::ExecAll()`. The code flow is illustrated in Figure 1. First, we explain the `Kernel::Kernel()` function, then detail how the components involved in initializing and executing a thread work, to provide an understanding of how the system runs a single program.

1.1 Kernel::Kernel()

`Kernel::Kernel()` (shown in Figure 2) initializes system settings, including console devices, network reliability, and command-line flags for debugging, and other configurations. If the `-e` flag is specified (line 49), the system stores the file name into `execfile`, allowing it to access the program file the user wants to execute.

1.2 Initialization

1.2.1 Kernel::ExecAll()

`Kernel::ExecAll()`, in Figure 3, is called by the initial thread to execute each user program by invoking `Kernel::Exec()` for each program file store in `execfile`. After completing all executions, the thread calls `Thread::Finish()` to trigger a context switch.

```

28 Kernel::Kernel(int argc, char **argv) {
29     randomSlice = FALSE;
30     debugUserProg = FALSE;
31     execExit = FALSE;
32     consoleIn = NULL; // default is stdin
33     consoleOut = NULL; // default is stdout
34     #ifndef FILESYS_STUB
35         formatFlag = FALSE;
36     #endif
37     reliability = 1; // network reliability, default is 1.0
38     hostName = 0; // machine id, also UNIX socket name
39                 // 0 is the default machine id
40     for (int i = 1; i < argc; i++) {
41         if (strcmp(argv[i], "-rs") == 0) {
42             ASSERT(i + 1 < argc);
43             RandomInit(atoi(argv[i + 1])); // initialize pseudo-random
44                                             // number generator
45             randomSlice = TRUE;
46             i++;
47         } else if (strcmp(argv[i], "-s") == 0) {
48             debugUserProg = TRUE;
49         } else if (strcmp(argv[i], "-e") == 0) {
50             execfile[++execfileNum] = argv[++i];

```

Figure 2: Kernel::Kernel()

```

void Kernel::ExecAll() {
    for (int i = 1; i <= execfileNum; i++) {
        int a = Exec(execfile[i]);
    }
    currentThread->Finish();
    // Kernel::Exec();
}

```

Figure 3: Kernel::ExecAll()

```

263 int Kernel::Exec(char *name) {
264     t[threadNum] = new Thread(name, threadNum);
265     t[threadNum]->setIsExec();
266     t[threadNum]->space = new AddrSpace();
267     t[threadNum]->Fork((VoidFunctionPtr)&ForkExecute, (void *)t[threadNum]);
268     threadNum++;
269
270     return threadNum - 1;
271
272     /*
273     cout << "Total threads number is " << execfileNum << endl;
274     for (int n=1;n<=execfileNum;n++) {
275         t[n] = new Thread(execfile[n]);
276         t[n]->space = new AddrSpace();
277         t[n]->Fork((VoidFunctionPtr) &ForkExecute, (void *)t[n]);
278         cout << "Thread " << execfile[n] << " is executing." << endl;
279     }
280     cout << "debug Kernel::Run finished.\n";
281     */
282     // Thread *t1 = new Thread(execfile[1]);
283     // Thread *t1 = new Thread("../test/test1");
284     // Thread *t2 = new Thread("../test/test2");
285
286     // AddrSpace *halt = new AddrSpace();
287     // t1->space = new AddrSpace();
288     // t2->space = new AddrSpace();
289
290     // halt->Execute("../test/halt");
291     // t1->Fork((VoidFunctionPtr) &ForkExecute, (void *)t1);
292     // t2->Fork((VoidFunctionPtr) &ForkExecute, (void *)t2);
293
294     // currentThread->Finish();
295     // Kernel::Run();
296     // cout << "after ThreadedKernel::Run();" << endl; // unreachable
297 }

```

Figure 4: Kernel::Exec()

1.2.2 Kernel::Exec()

`Kernel::Exec()` (shown in in Figure 4) creates a new thread, currently represented as a process, passes the function address and its arguments to `Fork()`, and initializes its thread (process) control block (a `Thread` object is a thread control block).

1.2.3 AddrSpace::AddrSpace()

```
65 AddrSpace::AddrSpace() {
66     pageTable = new TranslationEntry[NumPhysPages];
67     for (int i = 0; i < NumPhysPages; i++) {
68         pageTable[i].virtualPage = i; // for now, virt page # = phys page #
69         pageTable[i].physicalPage = i;
70         pageTable[i].valid = TRUE;
71         pageTable[i].use = FALSE;
72         pageTable[i].dirty = FALSE;
73         pageTable[i].readOnly = FALSE;
74     }
75
76     // zero out the entire address space
77     bzero(kernel->machine->mainMemory, MemorySize);
78 }
79
```

Figure 5: `AddrSpace::AddrSpace()`

`AddrSpace::AddrSpace()`, called by `Kernel::Exec()`, creates a page table for the thread and clears the entire main memory (see Figure 17).

The function uses `TranslationEntry` objects to map logical memory to physical memory. Since the system currently supports only uniprogramming, this mapping is one-to-one, allowing the entire main memory to be reset.

1.2.4 Thread::Fork()

```
91 void Thread::Fork(VoidFunctionPtr func, void *arg) {
92     Interrupt *interrupt = kernel->interrupt;
93     Scheduler *scheduler = kernel->scheduler;
94     IntStatus oldLevel;
95
96     DEBUG(dbgThread, "Forking thread: " << name << " f(a): " << (int)func << " " << arg);
97     StackAllocate(func, arg);
98
99     oldLevel = interrupt->SetLevel(IntOff);
100     scheduler->ReadyToRun(this); // ReadyToRun assumes that interrupts
101     | | | | | | | // are disabled!
102     (void)interrupt->SetLevel(oldLevel);
103 }
104
```

Figure 6: `Thread::Fork()`

`Thread::Fork()`, also called by `Kernel::Exec()`, allocates and initializes a stack for the thread using `StackAllocate()`. Finally, it puts the thread on the ready queue with interrupts disabled. The function is shown in Figure 6.

1.2.5 Thread::StackAllocate()

`Thread::StackAllocate()`, called by `Thread::Fork()`, initializes the stack for the thread (see Figure 7). The function sets the stack top based on the hardware architecture. It then places `ThreadRoot`, `ThreadBegin`, the function pointer and its arguments, and `ThreadFinish` on the stack. `ThreadRoot` enables interrupts, allowing the thread to execute upon context switching and ensures the thread stops upon completion. Specifically, `ThreadBegin` and `ThreadFinish` are invoked by `ThreadRoot` to call `Thread::Begin()` and `Thread::Finish()`, respectively.

1.2.6 Scheduler::ReadyToRun()

`Scheduler::ReadyToRun()` (see Figure 8), also called by `Thread::Fork()`, sets the thread status to ready and places it on the ready queue after ensuring interrupts are disabled.

1.2.7 Thread::Finish()

After placing all threads on the ready queue, the system returns to `Kernel::ExecAll()` and calls `Thread::Finish()`. `Thread::Finish()`, shown in Figure 9, first disables interrupts for `Sleep()`. If all threads have finished, it deallocates the thread; otherwise, it calls `Sleep()` to perform a context switch, as the thread may still be needed by other threads.

1.2.8 Thread::Sleep()

`Thread::Sleep()` (see Figure 10) blocks the current thread and attempts to retrieve the first thread from the ready queue. If no threads are available, the system enables interrupts to allow other programs to execute. Otherwise, it calls `Run()` to perform a context switch.

1.2.9 Scheduler::Run()

`Scheduler::Run()`, shown in Figure 11, is to perform a context switch. For each context switching, the system will save CPU state and machine state for the space for each user program. If the current thread has finished, it will be mark as destroyed and the system executes the next thread and invokes `Thread::Begin()`

```

301 void Thread::StackAllocate(VoidFunctionPtr func, void *arg) {
302     stack = (int *)AllocBoundedArray(StackSize * sizeof(int));
303
304     #ifdef PARISC
305         // HP stack works from low addresses to high addresses
306         // everyone else works the other way: from high addresses to low addresses
307         stackTop = stack + 16; // HP requires 64-byte frame marker
308         stack[StackSize - 1] = STACK_FENCEPOST;
309     #endif
310
311     #ifdef SPARC
312         stackTop = stack + StackSize - 96; // SPARC stack must contains at
313         // least 1 activation record
314         // to start with.
315         *stack = STACK_FENCEPOST;
316     #endif
317
318     #ifdef PowerPC // RS6000
319         stackTop = stack + StackSize - 16; // RS6000 requires 64-byte frame marker
320         *stack = STACK_FENCEPOST;
321     #endif
322
323     #ifdef DECMIPS
324         stackTop = stack + StackSize - 4; // -4 to be on the safe side!
325         *stack = STACK_FENCEPOST;
326     #endif
327
328     #ifdef ALPHA
329         stackTop = stack + StackSize - 8; // -8 to be on the safe side!
330         *stack = STACK_FENCEPOST;
331     #endif
332
333     #ifdef x86
334         // the x86 passes the return address on the stack. In order for SWITCH()
335         // to go to ThreadRoot when we switch to this thread, the return address
336         // used in SWITCH() must be the starting address of ThreadRoot.
337         stackTop = stack + StackSize - 4; // -4 to be on the safe side!
338         *(--stackTop) = (int)ThreadRoot;
339         *stack = STACK_FENCEPOST;
340     #endif
341
342     #ifdef PARISC
343         machineState[PCState] = PLabelToAddr(ThreadRoot);
344         machineState[StartupPCState] = PLabelToAddr(ThreadBegin);
345         machineState[InitialPCState] = PLabelToAddr(func);
346         machineState[InitialArgState] = arg;
347         machineState[WhenDonePCState] = PLabelToAddr(ThreadFinish);
348     #else
349         machineState[PCState] = (void *)ThreadRoot;
350         machineState[StartupPCState] = (void *)ThreadBegin;
351         machineState[InitialPCState] = (void *)func;
352         machineState[InitialArgState] = (void *)arg;
353         machineState[WhenDonePCState] = (void *)ThreadFinish;
354     #endif
355 }
356

```

Figure 7: Thread::StackAllocate()

```

55 void Scheduler::ReadyToRun(Thread *thread) {
56     ASSERT(kernel->interrupt->getLevel() == IntOff);
57     DEBUG(dbgThread, "Putting thread on ready list: " << thread->getName());
58     // cout << "Putting thread on ready list: " << thread->getName() << endl;
59     thread->setStatus(READY);
60     readyList->Append(thread);
61 }
62

```

Figure 8: Scheduler::ReadyToRun()


```

167 void Thread::Finish() {
168     (void)kernel->interrupt->SetLevel(IntOff);
169     ASSERT(this == kernel->currentThread);
170
171     DEBUG(dbgThread, "Finishing thread: " << name);
172     if (kernel->execExit && this->getIsExec()) {
173         kernel->execRunningNum--;
174         if (kernel->execRunningNum == 0) {
175             kernel->interrupt->Halt();
176         }
177     }
178     Sleep(TRUE); // invokes SWITCH
179     // not reached
180 }

```

Figure 9: Thread::Finish()

```

236 void Thread::Sleep(bool finishing) {
237     Thread *nextThread;
238
239     ASSERT(this == kernel->currentThread);
240     ASSERT(kernel->interrupt->getLevel() == IntOff);
241
242     DEBUG(dbgThread, "Sleeping thread: " << name);
243     DEBUG(dbgTraCode, "In Thread::Sleep, Sleeping thread: " << name << ", " << kernel->stats->totalTicks);
244
245     status = BLOCKED;
246     // cout << "debug Thread::Sleep " << name << "wait for Idle\n";
247     while ((nextThread = kernel->scheduler->FindNextToRun()) == NULL) {
248         kernel->interrupt->Idle(); // no one to run, wait for an interrupt
249     }
250     // returns when it's time for us to run
251     kernel->scheduler->Run(nextThread, finishing);
252 }
253

```

Figure 10: Thread::Sleep()

```

99 void Scheduler::Run(Thread *oThread, bool finishing) {
100     Thread *oldThread = kernel->currentThread;
101
102     ASSERT(kernel->interrupt->getLevel() == IntOff);
103
104     if (finishing) { // mark that we need to delete current thread
105         ASSERT(toBeDestroyed == NULL);
106         toBeDestroyed = oldThread;
107     }
108
109     if (oldThread->space != NULL) { // if this thread is a user program,
110         oldThread->SaveUserState(); // save the user's CPU registers
111         oldThread->space->SaveState();
112     }
113
114     oldThread->CheckOverflow(); // check if the old thread
115     | | | | | // had an undetected stack overflow
116
117     kernel->currentThread = nextThread; // switch to the next thread
118     nextThread->setStatus(RUNNING); // nextThread is now running
119
120     DEBUG(dbgThread, "Switching from: " << oldThread->getName() << " to: " << nextThread->getName());
121
122     // This is a machine-dependent assembly language routine defined
123     // in switch.s. You may have to think
124     // a bit to figure out what happens after this, both from the point
125     // of view of the thread and from the perspective of the "outside world".
126
127     SWITCH(oldThread, nextThread);
128
129     // we're back, running oldThread
130
131     // interrupts are off when we return from switch!
132     ASSERT(kernel->interrupt->getLevel() == IntOff);
133
134     DEBUG(dbgThread, "Now in thread: " << oldThread->getName());
135
136     CheckToBeDestroyed(); // check if thread we were running
137     | | | | | // before this one has finished
138     | | | | | // and needs to be cleaned up
139
140     if (oldThread->space != NULL) { // if there is an address space
141         oldThread->RestoreUserState(); // to restore, do it.
142         oldThread->space->RestoreState();
143     }
144 }
145

```

Figure 11: Scheduler::Run()

of the next thread. Otherwise, after context switching, this thread will restore its data by writing the register state into machine registers to continue executing.

In the initialization phase, once the main thread finishes, the system performs a context switch to the next thread, marking the end of initialization.

1.3 Executing a Thread

1.3.1 Thread::Begin()

```
141 void Thread::Begin() {
142     ASSERT(this == kernel->currentThread);
143     DEBUG(dbgThread, "Beginning thread: " << name);
144
145     kernel->scheduler->CheckToBeDestroyed();
146     kernel->interrupt->Enable();
147     if (kernel->execExit && this->getIsExec()) {
148         kernel->execRunningNum++;
149     }
150 }
```

Figure 12: Thread::Begin()

Thread::Begin() is called by ThreadRoot after a context switch (see Figure 12). The function first releases resources of the previous thread if it has finished by calling Scheduler::CheckToBeDestroyed(), then enables interrupts to support time-sharing. Finally, it updates the count of running threads in the system.

1.3.2 Scheduler::CheckToBeDestroyed()

Scheduler::CheckToBeDestroyed(), as Figure 13 shows, deletes deletes the finished thread by deallocating its stack.

1.3.3 ForkExecute()

ForkExecute(), shown in Figure 14, is used to execute forked threads. If the current thread's data has not been loaded, it calls Load(). The system then uses the thread's stack, which stores the functions and their arguments, to execute the program.

```

154 void Scheduler::CheckToBeDestroyed() {
155     if (toBeDestroyed != NULL) {
156         delete toBeDestroyed;
157         toBeDestroyed = NULL;
158     }
159 }

```

Figure 13: Scheduler::CheckToBeDestroyed()

```

247 void ForkExecute(Thread *t) {
248     if (!t->space->Load(t->getName())) {
249         return; // executable not found
250     }
251     t->space->Execute(t->getName());
252 }
253
254

```

Figure 14: ForkExecute()

1.3.4 AddrSpace::Load()

AddrSpace::Load(), shown in Figure 15, loads the code and data of a user program into main memory by first segmenting and then paging.

It opens the program file and reads a special header, `noffH`, to get information about the program's segments. Using this header, it calculates the required memory size and the number of pages. If there is enough physical memory, it loads the code and data segments into memory using virtual address, where virtual addresses currently match physical addresses. Finally, it closes the file and returns `TRUE` to indicate successful loading.

1.3.5 AddrSpace::Execute()

AddrSpace::Execute(), in Figure 16, initializes the machine by calling `InitRegisters()` to set up the machine registers, including the stack register, the initial program counter, and the next program counter. It also specifies the page table location for the process using `RestoreState()` (in contrast to context switching, where it retrieves the machine state).

After this setup, it calls `machine->Run()` to begin executing the program (handling instruction fetching, etc.).

Machine::Run() continuously fetches and executes instructions one by one. If resources are insufficient, the current thread goes to sleep until resources become available. Execution continues until an `SC_Exit()` exception occurs, which trig-

```

99 bool AddrSpace::Load(char *fileName) {
100     OpenFile *executable = kernel->fileSystem->Open(fileName);
101     NoFHHeader noFH;
102     unsigned int size;
103
104     if (executable == NULL) {
105         cerr << "Unable to open file " << fileName << ".ln";
106         return FALSE;
107     }
108
109     executable->ReadAt((char *)&noFH, sizeof(noFH), 0);
110     if ((noFH.noFhMagic != NOFPMAGIC) &&
111         (WordToHost(noFH.noFhMagic) == NOFPMAGIC))
112         SwapHeader(&noFH);
113     ASSERT(noFH.noFhMagic == NOFPMAGIC);
114
115 #ifdef RDATA
116     // how big is address space?
117     size = noFH.code.size + noFH.readonlyData.size + noFH.initData.size +
118         noFH.uninitData.size + UserStackSize;
119     // we need to increase the size
120     // to leave room for the stack
121 #else
122     // how big is address space?
123     size = noFH.code.size + noFH.initData.size + noFH.uninitData.size + UserStackSize; // we need to increase the size
124     // to leave room for the stack
125 #endif
126     numPages = divRoundUp(size, PageSize);
127     size = numPages * PageSize;
128
129     ASSERT(numPages <= NumPhysPages); // check we're not trying
130     // to run anything too big --
131     // at least until we have
132     // virtual memory
133
134     DEBUG(dbgAddr, "Initializing address space: " << numPages << ", " << size);
135
136     // then, copy in the code and data segments into memory
137     // Note: this code assumes that virtual address = physical address
138     if (noFH.code.size > 0) {
139         DEBUG(dbgAddr, "Initializing code segment.");
140         DEBUG(dbgAddr, noFH.code.virtualAddr << ", " << noFH.code.size);
141         executable->ReadAt(
142             &(kernel->machine->mainMemory[noFH.code.virtualAddr]),
143             noFH.code.size, noFH.code.inFileAddr);
144     }
145     if (noFH.initData.size > 0) {
146         DEBUG(dbgAddr, "Initializing data segment.");
147         DEBUG(dbgAddr, noFH.initData.virtualAddr << ", " << noFH.initData.size);
148         executable->ReadAt(
149             &(kernel->machine->mainMemory[noFH.initData.virtualAddr]),
150             noFH.initData.size, noFH.initData.inFileAddr);
151     }
152 #ifdef RDATA
153     if (noFH.readonlyData.size > 0) {
154         DEBUG(dbgAddr, "Initializing read only data segment.");
155         DEBUG(dbgAddr, noFH.readonlyData.virtualAddr << ", " << noFH.readonlyData.size);
156         executable->ReadAt(
157             &(kernel->machine->mainMemory[noFH.readonlyData.virtualAddr]),
158             noFH.readonlyData.size, noFH.readonlyData.inFileAddr);
159     }
160 #endif
161 #endif
162     delete executable; // close file
163     return TRUE; // success
164 }
165
166

```

Figure 15: AddrSpace::Load()

```

176 void AddrSpace::Execute(char *fileName) {
177     kernel->currentThread->space = this;
178
179     this->InitRegisters(); // set the initial register values
180     this->RestoreState(); // load page table register
181
182     kernel->machine->Run(); // jump to the user program
183
184     ASSERTNOTREACHED(); // machine->Run never returns;
185     // the address space exits
186     // by doing the syscall "exit"
187 }

```

Figure 16: AddrSpace::Execute()

gers `kernel->currentThread->Finish()`. The `Thread::Finish()` function has been explained above, so we will skip it here.

1.4 Questions in Spec

- **How does Nachos allocate the memory space for a new thread (process)?**

In `Kernel::Exec()`, a new `AddrSpace` is assigned to the new thread. The `AddrSpace` constructor initializes a page table using `TranslationEntry` to map logical memory directly to physical memory, as the system currently supports only uniprogramming.

- **How does Nachos initialize the memory content of a thread (process), including loading the user binary code in the memory?**

As mention, the `AddrSpace` constructor initializes a page table with `TranslationEntry` objects, mapping logical memory directly to physical memory. It sets the `valid` field to `TRUE` and the `use`, `dirty`, and `readOnly` bits to `FALSE`.

In `Kernel::Exec()`, `Thread::Fork()` is called with the function address and arguments. `Thread::Fork()` then calls `Thread::StackAllocate()`, which places the function pointer and arguments on the stack.

When `ForkExecute()` is invoked, it calls `AddrSpace::Load()` to load the user binary code. `AddrSpace::Load()` uses a special header (`noffH`) to load each segment, mapping virtual addresses to physical memory since they are currently identical.

- **How does Nachos create and manage the page table?**

As mention, in `Kernel::Exec()`, a new `AddrSpace` is created and stored in the `space` attribute of the new thread. The `AddrSpace` constructor initializes a page table using `TranslationEntry` to map logical memory directly to physical memory, as the system currently supports only uniprogramming. This setup allows the system to access the page table through `space`.

- **How does Nachos translate addresses?**

When loading a program file into main memory, `AddrSpace::Load()` loads the user binary code using a special header (`noffH`) by segmenting and paging as mention above. `AddrSpace::Load()` uses virtual addresses, which currently match physical addresses, to load each segment into memory.

During execution, `Machine::OneInstruction()` calls `Machine::ReadMem()`, which in turn calls `Machine::Translate()`. `Machine::Translate()` uses the virtual address and page size to determine the virtual page number and the offset within the page. It accesses the `physicalPage` attribute in the page table entry to get the frame number, allowing it to compute the physical address.

- **How does Nachos initialize the machine status (registers, etc) before running a thread (process)?**

When a thread is executed, `ForkExecute()` calls `AddrSpace::Execute()`. `AddrSpace::Execute()` initializes the machine by calling `InitRegisters()`, which sets up the machine registers, including the stack register, the initial program counter, and the next program counter.

- **Which object in Nachos acts the role of process control block?**

`Thread`. Currently, each thread represents a single process. Thread objects include the thread ID, thread status, stack, page table, and other PCB-related information.

- **When and how does a thread get added into the ReadyToRun queue of Nachos CPU scheduler?**

A thread is added to the ready queue when it is forked in `Thread::Fork()`. `Scheduler::ReadyToRun()`, called by `Thread::Fork()`, sets the thread status to ready and places it on the ready queue.

If a thread is waiting for resources, `Semaphore::V()` will call `Scheduler::ReadyToRun()` once the resources become available, adding the thread back to the ready queue.

2 Page Table Implementation

In order to support multiprogramming, we need to take the following steps:

1. Construct a page table with an appropriate number of entries for each process.
2. Check whether there are enough frames to accommodate the process. If so, assign the available frames to it.

3. Load the corresponding code into these frames.
4. Finally, if the process has completed its work and is to be terminated, delete the page table and return the frames to the kernel.

2.1 Construct a page table

For the following explanation, please refer to Figures 17 through 20.

Unlike the original implementation in NachOS, multiprogramming requires constructing the page table based on the size of each process. Therefore, we first comment out the code in `AddrSpace::AddrSpace()` (As shown in Figure 17) and create a page table constructor that takes `numPages` as an argument (As shown in Figure 18 and Figure 19), which will be invoked in `AddrSpace::Load()` (As shown in Figure 20, line 164.). (Since we need to know the size of the process to construct the page table, which is only available in `AddrSpace::Load()`.)

```

66  AddrSpace::AddrSpace() {
67      // pageTable = new TranslationEntry[NumPhysPages];
68      // for (int i = 0; i < NumPhysPages; i++) {
69          //     pageTable[i].virtualPage = i;    // for now, virt page # = phys page #
70          //     pageTable[i].physicalPage = i;    // modify in Load()
71          //     pageTable[i].valid = TRUE;
72          //     pageTable[i].use = FALSE;
73          //     pageTable[i].dirty = FALSE;
74          //     pageTable[i].readOnly = FALSE;
75          // }
76
77      // // zero out the entire address space
78      // bzero(kernel->machine->mainMemory, MemorySize);
79  }

```

Figure 17: Comment out `AddrSpace::AddrSpace()`

2.2 Assign the available frames

Before calling `AddrSpace::ConstructPageTable()`, we need to check the number of free frames to ensure that the process has enough space in memory. To track free frames, we add `int frameTable[NumPhysPages]` to record empty frames to the Kernel class in `kernel.h`. In this array, the index represents the frame number of a physical page, and if a frame is available, its value is set to 0 (and vice versa). Additionally, we add `int NumFreeFrame`, which represents the number of free frames in memory, too. Please refer to Figure 21 and Figure 22.

In addition, we add a new exception type, `MemoryLimitException`, in `machine.h` and set up a corresponding exception handler in `exception.cc` to handle situations where there is not enough space for the process (On line 153, Figure 20). You may refer to Figure 23 and Figure 24 for detail.


```

22 class AddrSpace {
23 public:
24     AddrSpace(); // Create an address space.
25     ~AddrSpace(); // De-allocate an address space
26
27     void ConstructPageTable(int numPages);
28
29     bool Load(char *fileName); // Load a program into addr space from
30     // a file
31     // return false if not found
32
33     void Execute(char *fileName); // Run a program
34     // assumes the program has already
35     // been loaded
36
37     void SaveState(); // Save/restore address space-specific
38     void RestoreState(); // info on a context switch
39
40     // Translate virtual address _vaddr_
41     // to physical address _paddr_. _mode_
42     // is 0 for Read, 1 for Write.
43     ExceptionType Translate(unsigned int vaddr, unsigned int *paddr, int mode);
44
45 private:
46     TranslationEntry *pageTable; // Assume linear page table translation
47     // for now!
48     unsigned int numPages; // Number of pages in the virtual
49     // address space
50
51     void InitRegisters(); // Initialize user-level CPU registers,
52     // before jumping to user code
53 };

```

Figure 18: AddrSpace::ConstructPageTable() Declaration

```

96 void AddrSpace::ConstructPageTable(int numPages) {
97     pageTable = new TranslationEntry[numPages];
98     for (int i = 0; i < numPages; i++) {
99         pageTable[i].virtualPage = i; // modify in Load()
100         pageTable[i].physicalPage = i; // modify in Load()
101         pageTable[i].valid = TRUE;
102         pageTable[i].use = FALSE;
103         pageTable[i].dirty = FALSE;
104         pageTable[i].readOnly = FALSE;
105     }
106 }

```

Figure 19: AddrSpace::ConstructPageTable()

```

118 bool AddrSpace::Load(char *fileName) {
119     OpenFile *executable = kernel->fileSystem->Open(fileName);
120     NoffHeader noffH;
121     unsigned int size;
122
123     if (executable == NULL) { ...
124
125     executable->ReadAt((char *)&noffH, sizeof(noffH), 0);
126     if ((noffH.noffMagic != NOFFMAGIC) &&
127         (WordToHost(noffH.noffMagic) == NOFFMAGIC))
128         SwapHeader(&noffH);
129     ASSERT(noffH.noffMagic == NOFFMAGIC);
130
131     #ifdef RDATA ...
132     #else ...
133     #endif
134
135     // DEBUG(dbgThread, "0x0 " << noffH.code.size << ", " << noffH.initData.size ...
136     numPages = divRoundUp(size, PageSize);
137     size = numPages * PageSize;
138
139     // Check #FreeFrames and #UsedPages
140     if (kernel->NumFreeFrame < numPages) ExceptionHandler(MemoryLimitException);
141
142     ASSERT(numPages <= NumPhysPages); // check we're not trying
143     // to run anything too big --
144     // at least until we have
145     // virtual memory
146
147     DEBUG(dbgThread,
148         | "Initializing address space: " << numPages << ", " << size);
149
150     // Construct page table then find empty frames to give to the process
151     kernel->currentThread->space->ConstructPageTable(numPages);
152     for (int i = 0, cnt = 0; i < NumPhysPages && cnt < numPages; i++) {
153         if (kernel->frameTable[i] == 0) {
154             kernel->frameTable[i] = 1;
155             kernel->NumFreeFrame--;
156         }
157     }
158 }

```

Figure 20: Invokes ConstructPageTable in AddrSpace::Load()

```

30 class Kernel {
31     int execRunningNum, // number of running threads
32
33     int hostName; // machine identifier
34     int frameTable[NumPhysPages];
35     int NumFreeFrame;
36
37     private:

```

Figure 21: int frameTable[NumPhysPages], int NumFreeFrame Declaration

```

28  Kernel::Kernel(int argc, char **argv) {
29      randomSlice = FALSE;
30      debugUserProg = FALSE;
31      execExit = FALSE;
32      consoleIn = NULL; // default is stdin
33      consoleOut = NULL; // default is stdout
34      frameTable[NumPhysPages] = {0};
35      NumFreeFrame = NumPhysPages;
36  #ifndef FILESYS_STUB
37      formatFlag = FALSE;
38  #endif
39      reliability = 1; // network reliability, default is

```

Figure 22: int frameTable[NumPhysPages], int NumFreeFrame Initialization

```

43  enum ExceptionType {
44      NoException,           // Everything ok!
45      SyscallException,      // A program executed a system call.
46      PageFaultException,    // No valid translation found
47      ReadOnlyException,     // Write attempted to page marked
48      | | | | |             // "read-only"
49      BusErrorException,     // Translation resulted in an
50      | | | | |             // invalid physical address
51      AddressErrorException, // Unaligned reference or one that
52      | | | | |             // was beyond the end of the
53      | | | | |             // address space
54      OverflowException,     // Integer overflow in add or sub.
55      IllegalInstrException, // Unimplemented or reserved instr.
56      MemoryLimitException,  // insufficient memory for a thread
57      NumExceptionTypes
58  };

```

Figure 23: New exception type: MemoryLimitException

```

50  void ExceptionHandler(ExceptionType which) {
51      char ch;
52      int val;
53      int type = kernel->machine->ReadRegister(2);
54      int status, exit, threadID, programID, fileID, numChar;
55      DEBUG(dbgSys, "Received Exception " << which << " type: " << type << "\n");
56      DEBUG(dbgTraCode, "In ExceptionHandler(), Received Exception "
57      | | | | |             << which << " type: " << type << ", "
58      | | | | |             << kernel->stats->totalTicks);
59      switch (which) {
60      case MemoryLimitException:
61          cerr << "Unexpected user mode exception " << (int)which << "\n";
62          ASSERT(false);

```

Figure 24: Handling MemoryLimitException in exception handler.

Finally, we identify available frames by iterating through `frameTable[NumPhysPages]` with a for loop and record them in the page table. As shown in Figure 25. It is worth noting that we should also clean up the frame assigned to this process to avoid using incorrect data (Figure 25, line 171.).

```

164     kernel->currentThread->space->ConstructPageTable(numPages);
165     for (int i = 0, cnt = 0; i < NumPhysPages && cnt < numPages; i++) {
166         if (kernel->frameTable[i] == 0) {
167             kernel->frameTable[i] = 1;
168             kernel->NumFreeFrame--;
169             kernel->currentThread->space->pageTable[cnt].physicalPage = i;
170             // zero out the entire address space
171             bzero(&kernel->machine->mainMemory[i * PageSize], PageSize);
172             cnt++;
173         }
174     }

```

Figure 25: Finding free frames withing `frameTable[NumPhysPages]`

2.3 Load the code

For those segments that are not read only, you may refer to Figure 26 and Figure 27 for more implementation details.

To place the code into the correct frame, we first determine its page number and offset, then calculate its physical address (You may refer to Figure 26, from line 179 to 184, and Figure 27, from line 196 to 202.). In this implementation, we use `ReadAt()` to read the code into memory. Here, the code is divided into code and `initData`, and the method for loading them are the same.

```

177     if (noffH.code.size > 0) {
178         // Calculate physical address to place code
179         int PhysAddr_base = kernel->currentThread->space
180             ->pageTable[noffH.code.virtualAddr / PageSize]
181             .physicalPage *
182             PageSize;
183         int PhysAddr_offset = noffH.code.virtualAddr % PageSize;
184         int physAddr = PhysAddr_base + PhysAddr_offset;
185
186         DEBUG(dbgThread, "Initializing code segment.");
187         DEBUG(dbgThread, "virtualAddr: " << noffH.code.virtualAddr
188             |<< " ", size: " << noffH.code.size);
189
190         executable->ReadAt(&(kernel->machine->mainMemory[physAddr]),
191             noffH.code.size, noffH.code.inFileAddr);
192         DEBUG(dbgThread, "Finish code segment initialization.");
193     }

```

Figure 26: Loading code into corresponding free frame.

For the segments that place read only data, we need to change the `readOnly` variable to "True". As shown in Figure 28.

```

194     if (noffH.initData.size > 0) {
195         // Calculate physical address to place initData
196         int PhysAddr_base =
197             kernel->currentThread->space
198             ->pageTable[noffH.initData.virtualAddr / PageSize]
199             .physicalPage *
200             PageSize;
201         int PhysAddr_offset = noffH.initData.virtualAddr % PageSize;
202         int physAddr = PhysAddr_base + PhysAddr_offset;
203
204         DEBUG(dbgThread, "Initializing data segment.");
205         DEBUG(dbgThread, "virtualAddr: " << noffH.initData.virtualAddr
206             | " | " << ", size: " << noffH.initData.size);
207         executable->ReadAt(&(kernel->machine->mainMemory[physAddr]),
208             | " | " << noffH.initData.size, noffH.initData.inFileAddr);
209         // May overwrite
210         DEBUG(dbgThread, "Finish code data initialization.");
211     }

```

Figure 27: Loading initial data into corresponding free frame.

```

213 #ifdef RDATA
214     if (noffH.readonlyData.size > 0) {
215         int PhysAddr_base =
216             kernel->currentThread->space
217             ->pageTable[noffH.readonlyData.virtualAddr / PageSize]
218             .physicalPage *
219             PageSize;
220         int PhysAddr_offset = noffH.readonlyData.virtualAddr % PageSize;
221         int physAddr = PhysAddr_base + PhysAddr_offset;
222
223         kernel->currentThread->space
224         ->pageTable[noffH.readonlyData.virtualAddr / PageSize]
225         .readOnly = TRUE;
226
227         DEBUG(dbgAddr, "Initializing read only data segment.");
228         DEBUG(dbgThread,
229             | "virtualAddr: " << noffH.readonlyData.virtualAddr
230             | " | " << ", size: " << noffH.readonlyData.size);
231         executable->ReadAt(&(kernel->machine->mainMemory[physAddr]),
232             | " | " << noffH.readonlyData.size,
233             | " | " << noffH.readonlyData.inFileAddr);
234     }

```

Figure 28: Loading read only data into corresponding free frame.

2.4 Return the frames

Once the process completes its work, we need to delete the page table. There are a few things to keep in mind. First, `frameTable[NumPhysPages]` needs to be updated back to 0. Second, `NumFreeFrame` should be incremented to reflect the freed frames. Figure 29 shows our implementation of `AddrSpace::~AddrSpace()`.

```
86 AddrSpace::~AddrSpace() {
87     for (int i = 0; i < numPages; ++i) {
88         kernel->frameTable[kernel->currentThread->space->pageTable[i]
89         | | | | |.physicalPage] = 0;
90         kernel->NumFreeFrame++;
91     }
92
93     delete pageTable;
94 }
```

Figure 29: Returning frames to the kernel.